Planning and Reinforcement Learning in AI Systems

(PARLAI)

Lesson 3a - Decision Making as Supervised Learning (part 1)

Sarah Keren

The Taub Faculty of Computer Science
Technion - Israel Institute of Technology

So far: We explored various approaches to decision-making.

So far: We explored various approaches to decision-making.



So far: We explored various approaches to decision-making.



Today: focus on model-free supervised learning approaches.

Supervised Learning

Learning from Data

A central goal of **machine learning** is to automatically discover patterns from data. Instead of manually programming rules, we provide the system with data and allow it to learn a model.

Examples of tasks that can be learned from data:

- Recognizing objects in images
- Predicting house prices
- Detecting anomalies in sensor readings
- Classifying medical images



Different learning paradigms exist depending on the type of data and feedback available.

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Supervised Learning

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Supervised learning is a paradigm in machine learning where a model is trained using labeled examples.

It is one of the most common learning paradigms (and the core of most foundation models).

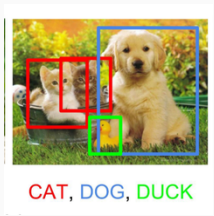
Supervised learning answers the question:

Assuming an unknown true distribution $\mathcal{D}$, how can we learn to make predictions from labeled examples sampled from $\mathcal{D}$?

Supervised Learning Examples

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

- Recognizing objects in images **using images in which objects were identified**.
- Predicting house prices **based on examples of house prices under certain conditions**
- Detecting anomalies in sensor readings **based on reading for which the actual values are known**
- Classifying medical images **based on labels provided by human doctors**

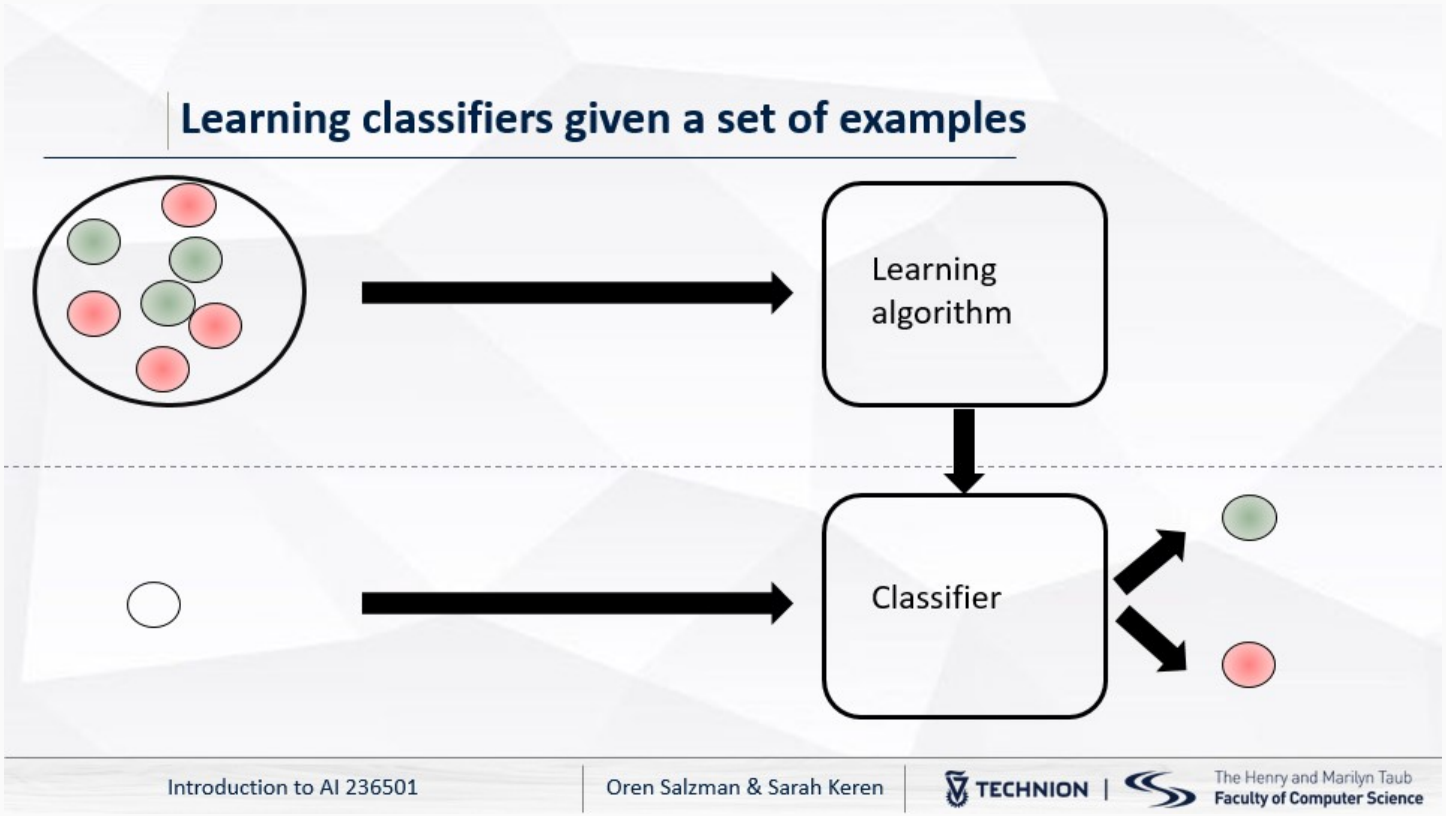








Supervised Learning in IAI



In IAI, we used supervised learning to learn the underlying model of the environment.
How is this related to computing policies and to decision-making?

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Algorithmic Setup

Supervised Learning Components

A distribution $\mathcal{D}$, from which samples are drawn i.i.d. (independent and identically distributed)

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Supervised Learning Components

A distribution $\mathcal{D}$, from which samples are drawn i.i.d. (independent and identically distributed)

A sample, or labeled example (x_i, y_i) is a pair consisting of:

- an **input** (feature vector) x_i
- the corresponding label or target value y_i

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Supervised Learning Components

A **distribution** $\mathcal{D}$, from which samples are drawn i.i.d. (independent and identically distributed)

A **sample**, or labeled example (x_i, y_i) is a pair consisting of:

- an **input** (feature vector) x_i
- the corresponding **label** or target value y_i

A **sample set**: $S = \{(x_i, y_i)\}_{i=1}^{m}$ is used to learn a **predictor** $h(x) \in H$ from a **hypothesis class** H that generalizes well with respect to the underlying distribution $\mathcal{D}$, i.e., predicts the correct label for new inputs $(x, y) \sim \mathcal{D}$ with high probability.

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Supervised Learning Components

A **distribution** $\mathcal{D}$, from which samples are drawn i.i.d. (independent and identically distributed)

A **sample**, or labeled example (x_i, y_i) is a pair consisting of:

- an **input** (feature vector) x_i
- the corresponding **label** or target value y_i

A **sample set**: $S = \{(x_i, y_i)\}_{i=1}^{m}$ is used to learn a **predictor** $h(x) \in H$ from a **hypothesis class** H that generalizes well with respect to the underlying distribution $\mathcal{D}$, i.e., predicts the correct label for new inputs $(x, y) \sim \mathcal{D}$ with high probability.

When the label y takes values in a **discrete set** (e.g., $\{0, 1\}$ or a set of classes), the task is called **classification**.

When the label y takes values in a **continuous space** (e.g., $y \in \mathbb{R}$), the task is called **regression**.

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Supervised Learning Setup: Example

A sample includes information x_i about a person (age, gender, city of birth, ...) and her height y_i.

$x_1 = \{35, F, NY\}, \quad y_1 = 1.58$

$x_2 = \{5, M, TLV\}, \quad y_2 = 0.7$

We seek a predictor $h(x) \in H$ among a hypothesis class H, e.g., linear functions, decision trees, neural networks.

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Supervised Learning Setup: Example

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

A sample includes information x_i about a person (age, gender, city of birth, ...) and her height y_i.

$$x_1 = \{35, F, NY\}, \quad y_1 = 1.58$$

$$x_2 = \{5, M, TLV\}, \quad y_2 = 0.7$$

We seek a predictor $h(x) \in H$ among a hypothesis class H, e.g., linear functions, decision trees, neural networks. For example, we may seek a linear model

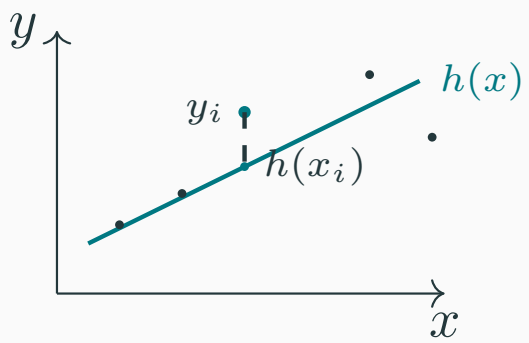
$$H = \left\{ h_w(x) = w^\top x + b \;\middle|\; w \in \mathbb{R}^d,\ b \in \mathbb{R} \right\}$$

where w are the weights and b is the bias.

Loss and Risk

In supervised learning, we aim to **learn a hypothesis that minimizes prediction error on unseen data**.

Loss $\ell(h(x), y)$: point-wise error on a **single example**



Expected Risk $\mathcal{R}(h)$: expected loss over the **data distribution**

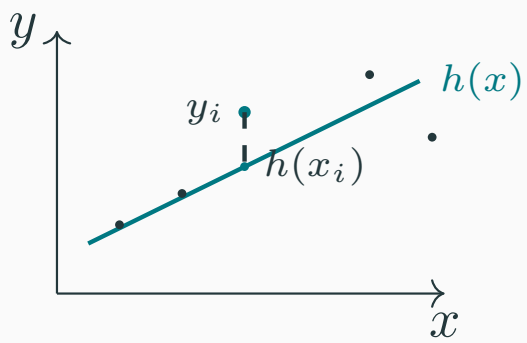
$$\mathcal{R}(h) = \mathbb{E}_{(x,y)\sim\mathcal{D}}[\ell(h(x), y)]$$

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Loss and Risk

In supervised learning, we aim to **learn a hypothesis that minimizes prediction error on unseen data**.

Loss $\ell(h(x), y)$: point-wise error on a **single example**



Expected Risk $\mathcal{R}(h)$: expected loss over the **data distribution**

$$\mathcal{R}(h) = \mathbb{E}_{(x,y)\sim\mathcal{D}}[\ell(h(x), y)]$$

Since we typically do not have access to the true distribution $\mathcal{D}$, we rely instead on a **finite dataset** $(x, y) \sim \mathcal{D}$ to compute the **Empirical Risk** $\hat{\mathcal{R}}(h)$, i.e., the average loss over the **training data**.

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Empirical Risk Minimization

Given a sample set $S = \{(x_i, y_i)\}_{i=1}^m$ and a hypothesis class H, we evaluate each hypothesis by its average loss on the training data.

For a hypothesis $h \in H$, the **empirical risk** is

$$\hat{\mathcal{R}}(h) = \frac{1}{m} \sum_{i=1}^{m} \ell(h(x_i), y_i)$$

The learning problem is to find a hypothesis with minimal empirical risk:

$$\hat{h} = \arg\min_{h \in H} \hat{\mathcal{R}}(h) = \arg\min_{h \in H} \frac{1}{m} \sum_{i=1}^{m} \ell(h \in h(x_i), y_i)$$

Interpretation: choose the predictor whose predictions incur the smallest average loss on the observed examples.

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Examples of Loss Functions

Classification

- 0-1 Loss

$$\ell(h(x), y) = \begin{cases} 0 & h(x) = y \\ 1 & h(x) \neq y \end{cases}$$

- Cross-Entropy Loss (used with probabilistic classifiers)

$$\ell(h(x), y) = -\log P(y|x)$$

Regression

- Squared Loss (Mean Squared Error)

$$\ell(h(x), y) = (h(x) - y)^2$$

- Absolute Loss

$$\ell(h(x), y) = |h(x) - y|$$

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Supervised Learning Setup

Input: A finite data set generated from an unknown distribution $(x, y) \sim \mathcal{D}$ over inputs and labels, and a hypothesis class H

The data set is assumed to be sampled **independently and identically distributed** from some underlying distribution $\mathcal{D}$.

Optimality: a hypothesis that minimizes the expected loss:

$$h^\star = \arg\min_{h \in H} \mathbb{E}_{(x,y)\sim\mathcal{D}} \left[\ell(h(x), y)\right]$$

Empirical Objective: Since the true distribution $\mathcal{D}$ is unknown, we approximate this objective using the empirical loss:

$$\hat{h} = \arg\min_{h \in H} \frac{1}{m} \sum_{i=1}^{m} \ell(h(x_i), y_i)$$

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Supervised Learning Setup

In practice, we split the available dataset into three parts:

Training set S_{train}:

- used to **learn the model parameters**
- optimize empirical risk:

Validation (evaluation) set S_{val}:

- used to **tune hyperparameters** (e.g., model size)
- helps detect **overfitting**

Test set S_{test}:

- used **only once at the end**
- provides an **unbiased estimate of generalization performance**

Never use test data during training or model selection.

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Algorithms

Supervised Learning Algorithms

There are several approaches to finding (or approximating) $\hat{h}$:

- **Analytical solutions** Closed-form minimization (e.g., linear regression with squared loss)
- **Iterative optimization** Gradient-based methods (e.g., SGD) for optimizing differentiable models, and search-based or gradient-free methods for non-differentiable models
- **Structured / symbolic methods** Search over hypothesis space (e.g., decision trees, rule learning)
- **Kernel / instance-based methods** Implicit hypothesis spaces (e.g., SVMs, nearest neighbors)
- **Pretrained / foundation models** Start from a large pretrained model and adapt:
  - fine-tuning on S
  - prompting / in-context learning

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

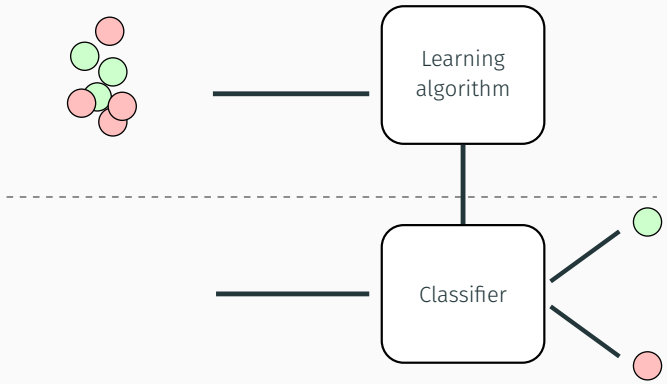
Example Approach: Iterative Supervised Learning via ERM

Iterative Supervised Learning via ERM

Require: Training set $S = \{(x_i, y_i)\}_{i=1}^{m}$, hypothesis class H, loss ℓ

Ensure: Learned hypothesis $\hat{h}$

1: Initialize hypothesis $h \in H$
2: **repeat**
3: Update h using optimization (e.g., gradient descent)
4: **until** convergence
5: $\hat{h} \leftarrow h$
6: **Return** $\hat{h}$



Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Example Update Rule: Gradient Descent

A common choice for the update function is **gradient descent**, which improves the current hypothesis by moving it in the direction that reduces the empirical risk (loss).

If we let the empirical risk be $\hat{\mathcal{R}}(h) = \frac{1}{m}\sum_{i=1}^{m} \ell(h(x_i), y_i)$, each iteration becomes:

$$h \leftarrow h - \eta \nabla \left(\frac{1}{m} \sum_{i=1}^{m} \ell\big(h(x_i), y_i\big) \right)$$

where:

- $\eta > 0$ is the **learning rate**
- $\nabla \hat{\mathcal{R}}(h)$ is the gradient of the empirical risk with respect to the model parameters

Gradient descent repeatedly adjusts the model parameters to reduce the training error.

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Example Update Rule: Gradient Descent

In modern methods, deep learning is often used.
We parameterize the hypothesis as $h_\theta \in H$ with parameters θ.
The empirical risk is: $\hat{\mathcal{R}}(\theta) = \frac{1}{m}\sum_{i=1}^{m} \ell\big(h_\theta(x_i), y_i\big)$ and each iteration is:

$$\theta \leftarrow \theta - \eta \nabla_\theta \left(\frac{1}{m} \sum_{i=1}^{m} \ell\big(h_\theta(x_i), y_i\big) \right)$$

.

We iteratively update the parameters to reduce the empirical risk until convergence.

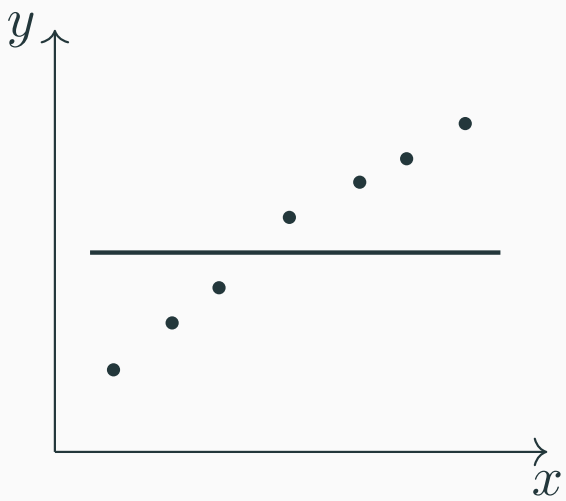
Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Challenges

Supervised Learning : Key challenges

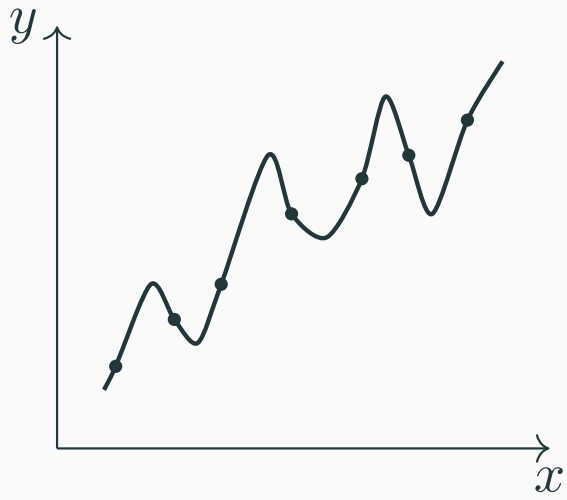
Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Underfitting



Avoid **underfitting**: using a model that is too simple to capture patterns.

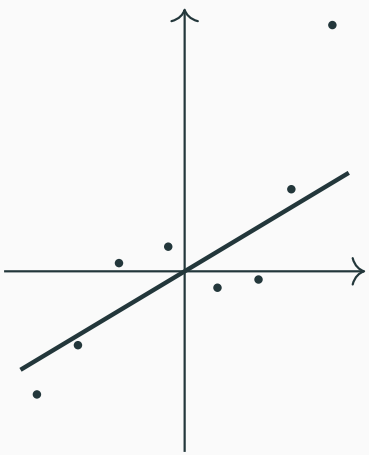
Overfitting



Avoid **overfitting**: fitting the training data too closely while failing to generalize.

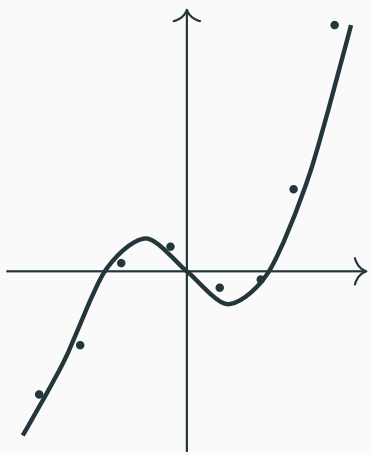
Example: Choosing Model Complexity

Assume the data is generated from an unknown cubic relationship: $y = x^3 - x + \epsilon$ where ϵ is a small noise term.



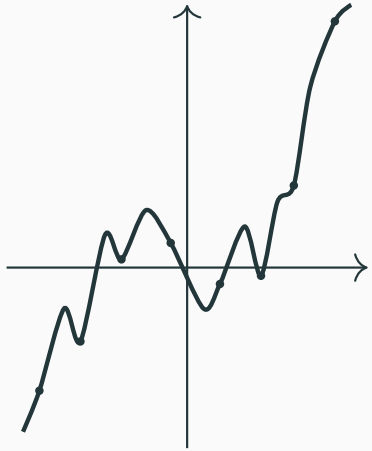
Fit with a **linear model**:

Too simple to capture the cubic trend.



Fit with a **degree-3 polynomial**:

Captures the underlying structure well.



Fit with a **n-degree polynomial**:

Fits the noise rather than the true pattern.

The goal is to find a model complexity that **generalizes well**.

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Finding an Effective Loss Function

A sample includes information x_i about a person (age, gender, city of birth, ...) and her height y_i.

$x_1 = \{35, F, NY\}, \quad y_1 = 1.58$

$x_2 = \{5, M, TLV\}, \quad y_2 = 0.7$

What is a helpful loss measure?

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Finding an Effective Loss Function

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM



What is a helpful loss measure?

Finding an Effective Loss Function

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

For some prediction tasks, there is no obvious notion of **distance between outputs**.

Possible choices encode different assumptions:

- **0–1 loss**: treat all incorrect predictions equally.
- **Cost matrix**: specify a penalty for each type of mistake.
- **Representation distance**: compare output embeddings, assuming their geometry reflects relevant similarity.

Challenge: How to define a loss that captures meaningful errors when the output space provides no natural distance?

Finding an Effective Loss Function

The loss $\ell(h(x), y)$ specifies how undesirable a prediction is.

Choosing it is part of **defining the learning problem**.

- **Meaning of the output:** Numerical differences are meaningful for heights, but not for arbitrary class identifiers.
- **Severity of errors:** Should a large error receive a proportionally larger penalty, or a much larger one?
- **Consequences of errors:** Are underestimation and overestimation equally costly?

A loss can express a notion of "distance,"
but need not be a mathematical metric.

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

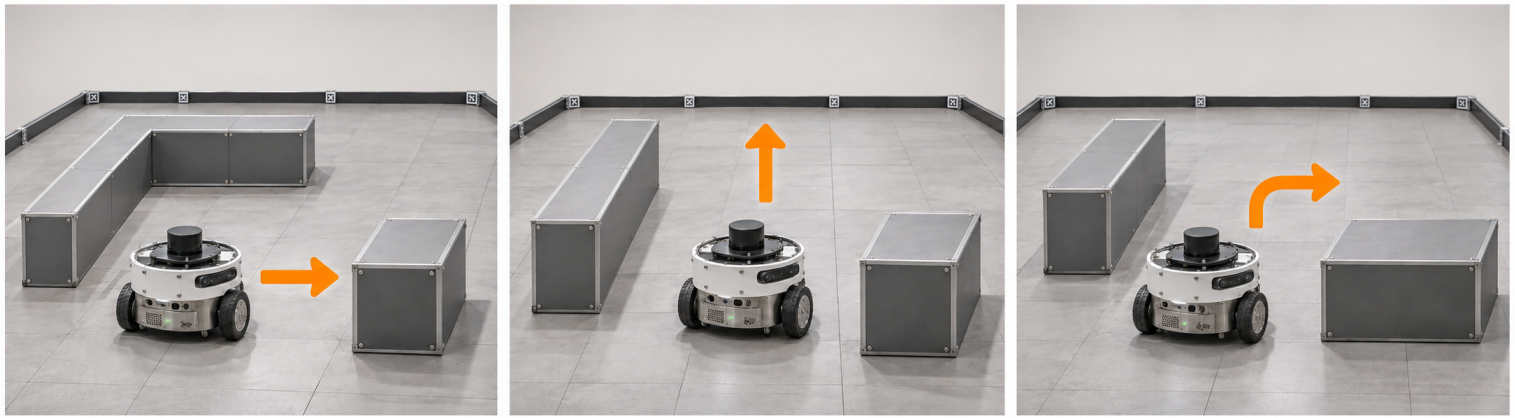
From SL to DM

From Supervised Learning to Decision-Making

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

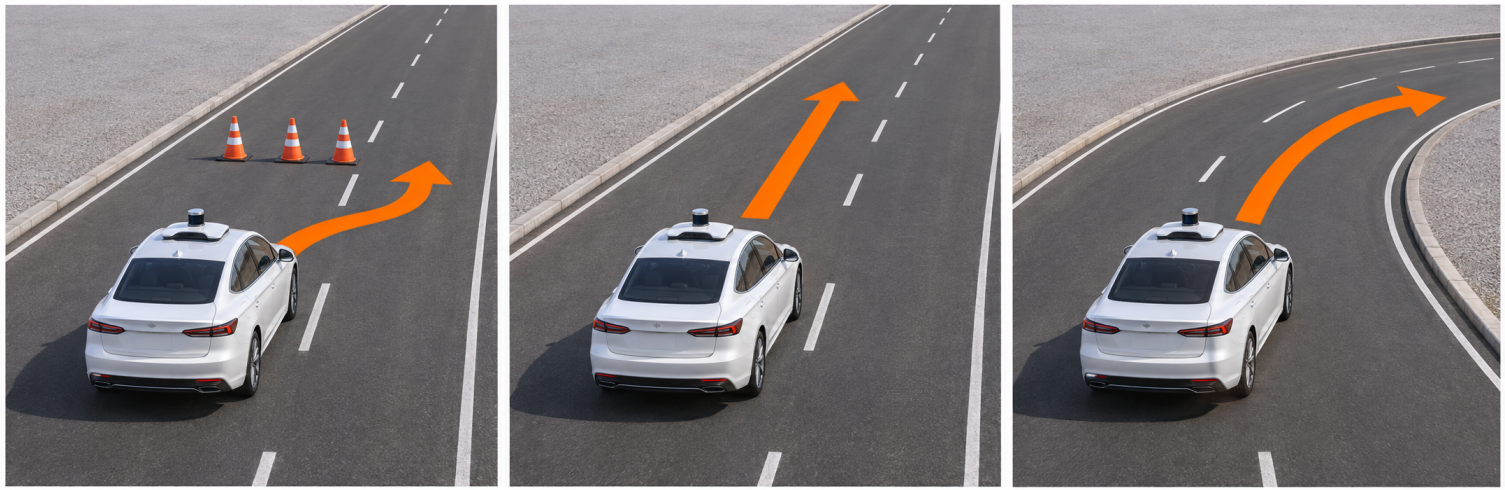
How does supervised learning relate to decision-making?

From Supervised Learning to Decision-Making



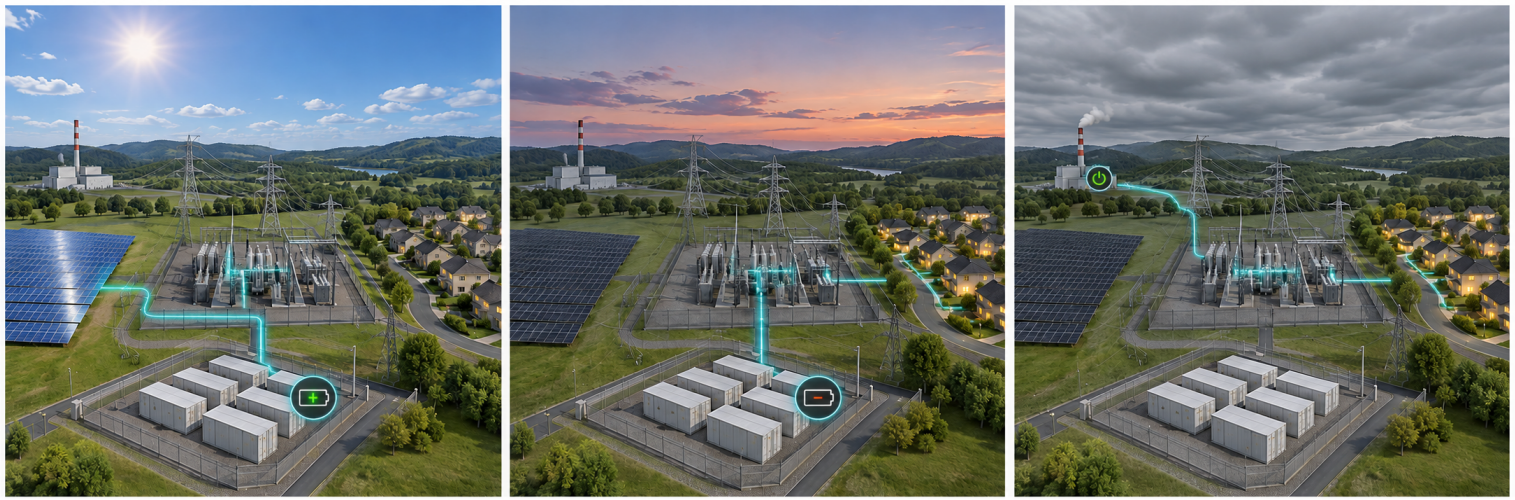
Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

From Supervised Learning to Decision-Making



Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

From Supervised Learning to Decision-Making



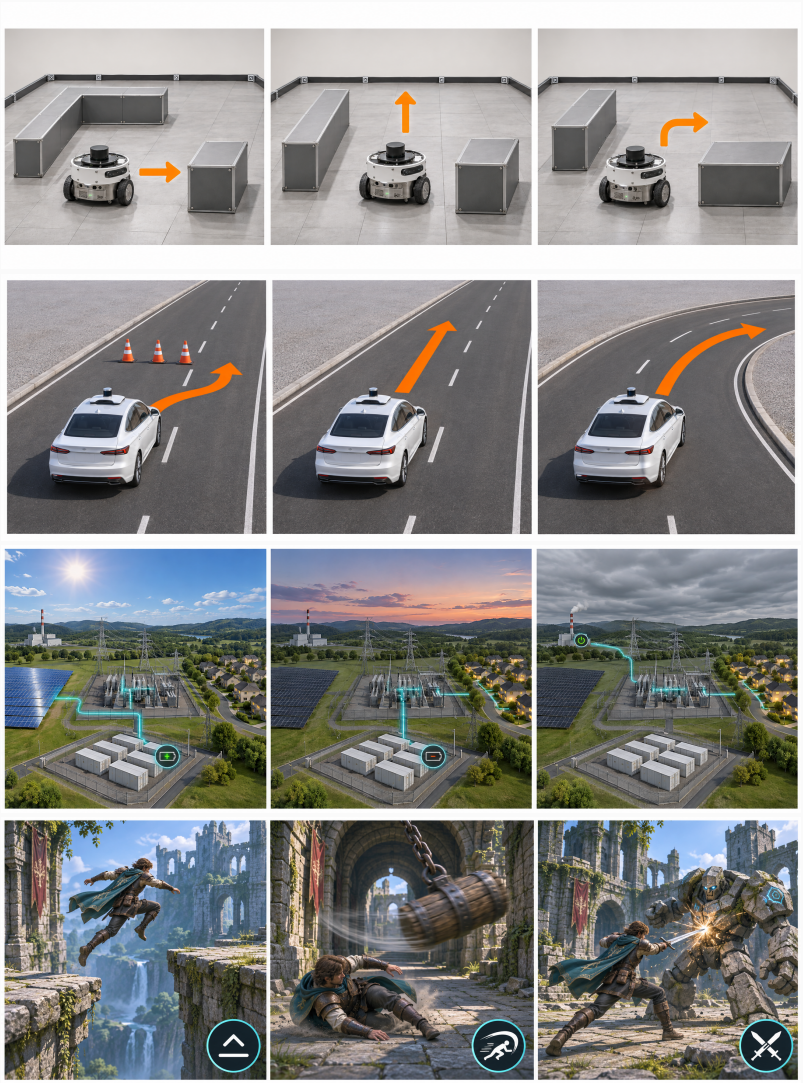
Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

From Supervised Learning to Decision-Making



Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

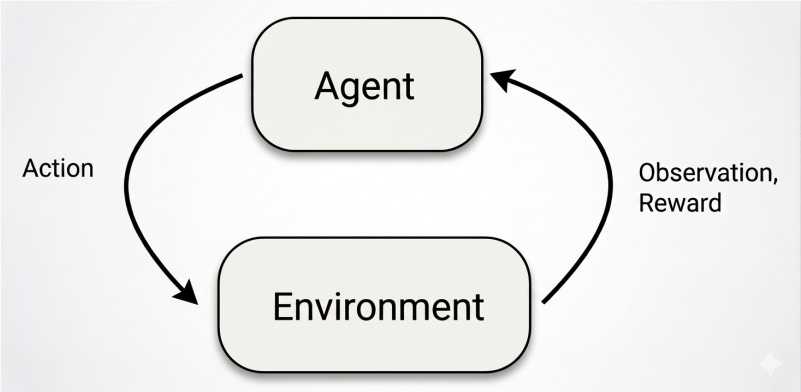
Challenges in supervised learning to decision-making?



From Supervised Learning to Decision-Making

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Supervised learning relies on the assumption that there is a given dataset and that data is **i.i.d.** from some underlying distribution $\mathcal{D}$.



Decision-making involves **interaction** with the environment.

From Supervised Learning to Decision-Making

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Challenge: Data is no longer i.i.d.

During interaction, observations are **temporally dependent**, and the agent's actions influence the data it encounters:

$$a_t \sim \pi(\cdot \mid s_t), \qquad s_{t+1} \sim P(\cdot \mid s_t, a_t).$$

Why does this matter?

From Supervised Learning to Decision-Making

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Challenge: Data is no longer i.i.d.

During interaction, observations are **temporally dependent**, and the agent's actions influence the data it encounters:

$$a_t \sim \pi(\cdot \mid s_t), \qquad s_{t+1} \sim P(\cdot \mid s_t, a_t).$$

Why does this matter?

- Consecutive samples are correlated, so additional samples may provide less new information than independent samples.
- Changing the policy changes the distribution of visited states - standard i.i.d. generalization guarantees do not directly apply:

From Supervised Learning to Decision-Making

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

Challenge: Distribution-specific guarantees

Even if we learn a model with low generalization error, the guarantees hold **only for the distribution** $\mathcal{D}$ from which the data was sampled.

When predictions are used for **decision-making** or **interaction with an environment**, this becomes critical:

- The data distribution during deployment may differ from $\mathcal{D}$
- There are **no guarantees** that performance will remain good under such shifts

From Supervised Learning to Decision-Making

Supervised Learning
Algorithmic Setup
Algorithms
Challenges
From SL to DM

How does supervised learning relate to decision-making?

We will explore two key paradigms:

- **Behavior cloning (imitation learning):** learn a policy from expert demonstrations, treating actions as labels.
- **Active learning:** select which data to acquire.